From the Bank Dataset, the columns we are dropping are ‘duration’, ‘default’ and ‘pdays’.

Duration of a call is not known beforehand and once we know the duration; we also know if the final answer is yes or no. Hence, we drop the duration column.
We drop the default column because it has too many 'unknown' values. We drop the 'pdays' column because it has high standard deviation and as less than 5% of the people have been contacted before.

Then, we change the ordinal values manually into numerical values ranging from -1 to 6 while preserving the order. We also changed the remaining categorical values to numerical values using get_dummies (which implements One-Hot Encoding). Then we split the data into training and test data.

Then we fit the following four models on both the datasets:

i) Decision Tree Classifier
ii) Naïve Bayes Classifier
iii) Random Forest Classifier
iv) Boosting on Decision Stumps

	Decision Tree Classifier	Naïve Bayes Classifier	Random Forest Classifier	Adaboost on Decision Stumps
Accuracy	83.44%	86.2%	83.09%	81.71%
Precision	35.47%	39.69%	35.17%	33.44%
Recall	58.46%	44.75%	60.45%	63.92%
F1 Score	44.15	42.07	44.47	43.91

Time and Memory consumed

	Decision Tree Classifier	Naïve Bayes Classifier (Gaussian)	Random Forest Classifier	Adaboost on Decision Stumps
Time	2.07 s	1.53 s	1.69 s	4.63 s
Peak Memory Usage	38.84 MB	27.49 MB	14.83 MB	17.03 MB